

INTRODUCTION AND MOTIVATION

Christian Kaestner

LECTURE LOGISTICS DURING A PANDEMIC

If you can hear me, open the participant panel in Zoom and check "yes"



SIMULATING IN-CLASS EXPERIENCE

Discussions and interactions are important. We'll have regular in-class discussions and exercises

- Use chat or "raise hand" feature
- Always keep camera on, muted by default
- Set preferred name in Zoom
- Attend lecture and recitation live, recordings only as backup
- I may call on you

- Suggestion: Have chat and participant list open, maybe separate window for gallery view, second monitor highly recommended
- **Contact me for accommodations!**

LEARNING GOALS

- Understand how AI components are parts of larger systems
- Illustrate the challenges in engineering an AI-enabled system beyond accuracy
- Explain the role of specifications and their lack in machine learning and the relationship to deductive and inductive reasoning
- Summarize the respective goals and challenges of software engineers vs data scientists

DISCLAIMER: EXPERIMENTAL CLASS

Second offering, but significant redesign



A Venn diagram consisting of two overlapping circles. The left circle is light green and contains the text 'Data Scientists'. The right circle is light orange and contains the text 'Software Engineers'. The overlapping area in the center is a darker shade of orange.

**Data
Scientists**

**Software
Engineers**

AGENDA



CASE STUDY: THE TRANSCRIPTION SERVICE STARTUP



GoTranscript education discount



Place Your Order



Login

Sign Up ✓



Contact us



▼ Services

Cost Estimate

▼ Samples

Pricing

About Us

Transcriptions samples

Captions and Subtitles samples

Academic Transcription Services

Our education transcription services have got you covered:

✓ Lectures

✓ Seminars

✓ Group discussions

✓ Interviews

✓ Presentations

20% discount for:



Chat with us

TRANSCRIPTION SERVICES

- Take audio or video files and produce text.
 - Used by academics to analyze interview text
 - Podcast show notes
 - Subtitles for videos
- State of the art: Manual transcription, often mechanical turk (1.5 \$/min)

THE STARTUP IDEA

PhD research on domain-specific speech recognition, that can detect technical jargon

DNN trained on public PBS interviews + transfer learning on smaller manually annotated domain-specific corpus

Research has shown amazing accuracy for talks in medicine, poverty and inequality research, and talks at Ruby programming conferences; published at top conferences

Idea: Let's commercialize the software and sell to academics and conference organizers

LIKELY CHALLENGES?



Everybody type 2 likely challenges in the chat but *do not send them yet*. Vote "yes" when done.

Speaker notes

Ask students to write at least 3 challenges each, collect answers on board



A Venn diagram consisting of two overlapping circles. The left circle is light green and contains the text 'Data Scientists'. The right circle is light orange and contains the text 'Software Engineers'. The overlapping area in the center is a darker shade of orange.

**Data
Scientists**

**Software
Engineers**

SOFTWARE ENGINEER

DATA SCIENTIST

- Often fixed dataset for training and evaluation (e.g., PBS interviews)
- Focused on accuracy
- Prototyping, often Jupyter notebooks or similar
- Expert in modeling techniques and feature engineering
- Model size, updateability, implementation stability typically does not matter

- Builds a product
- Concerned about cost, performance, stability, release time
- Identify quality through customer satisfaction
- Must scale solution, handle large amounts of data
- Detect and handle mistakes, preferably automatically
- Maintain, evolve, and extend the product over long periods
- Consider requirements for security, safety, fairness

LIKELY COLLABORATION CHALLENGES?



QUALITIES OF INTEREST ("ILITIES")

- Quality is about more than the absence of defects
- Quality in use (effectiveness, efficiency, satisfaction, freedom of risk, ...)
- Product quality (functional correctness and completeness, performance efficiency, compatibility, usability, dependability, scalability, security, maintainability, portability, ...)
- Process quality (manageability, evolvability, predictability, ...)
- "Quality is never an accident; it is always the result of high intention, sincere effort, intelligent direction and skillful execution; it represents the wise choice of many alternatives." (many attributions)

GARVIN'S EIGHT CATEGORIES OF PRODUCT QUALITY

- Performance
- Features
- Reliability
- Conformance
- Durability
- Serviceability
- Aesthetics
- Perceived Quality

Reference: Garvin, David A., [What Does Product Quality Really Mean](#). Sloan management review 25 (1984).

RELEVANT QUALITIES FOR TRANSCRIPTION SERVICE?



EXAMPLES FOR DISCUSSION

- What does correctness or accuracy really mean? What accuracy do customers care about?
 - How can we see how well we are doing in practice? How much feedback are customers going to give us before they leave?
 - Can we estimate how good our transcriptions are? How are we doing for different customers or different topics?
 - How to present results to the customers (including confidence)?
 - When customers complain about poor transcriptions, how to prioritize and what to do?
-
- What are unacceptable mistakes and how can they be avoided? Is there a safety risk?
 - Can we cope with an influx of customers?
 - Will transcribing the same audio twice produce the same result? Does it matter?
 - How can we debug and fix problems? How quickly?

EXAMPLES FOR DISCUSSION 2

- With more customers, transcriptions are taking longer and longer -- what can we do?
 - Transcriptions sometimes crash. What to do?
 - How do we achieve high availability?
 - How can we see that everything is going fine and page somebody if it is not?
 - We improve our entity detection model but somehow system behavior degrades... Why?
 - Tensorflow update; does our infrastructure still work?
 - Once somewhat successful, how to handle large amounts of data per day?
 - Buy more machines or move to the cloud?
-
- Models are continuously improved. When to deploy? Can we roll back?
 - Can we offer live transcription as an app? As a web service?
 - Can we get better the longer a person talks? Should we then go back and reanalyze the beginning? Will this benefit the next upload as well?

EXAMPLES FOR DISCUSSION 3

- How many domains can be supported? Do we have the server capacity?
- How specific should domains be? Medical vs "International Conference on Allergy & Immunology"?
- How to make it easy to support new domains?
- Can we handle accents?
- Better recognition of male than female speakers?
- Can and should we learn from customer data?
- How can we debug problems on audio files we are not allowed to see?
- Any chance we might private leak customer data?
- Can competitors or bad actors attack our system?

the-changelog-318

← [Dashboard](#)

Quality: High ⓘ

Last saved a few seconds ago

...

Share

00:00  Offset 00:00 01:31:27



Play



Back 5s

1x

Speed



Volume

NOTES

Write your notes here

Speaker 5 ▶ 07:44

Yeah. So there's a slight story behind that. So back when I was in, uh, Undergrad, I wrote a program for myself to measure a, the amount of time I did data entry from my father's business and I was on windows at the time and there wasn't a function called time dot [inaudible] time, uh, which I needed to parse dates to get back to time, top of representation, uh, I figured out a way to do it and I gave it to what's called the python cookbook because it just seemed like something other people could use. So it was just trying to be helpful. Uh, subsequently I had to figure out how to make it work because I didn't really have to. Basically, it bothered me that you had to input all the locale information and I figured out how to do it over the subsequent months. And actually as a graduation gift from my Undergrad, the week following, I solved it and wrote it all out.

Speaker 5 ▶ 08:38

And I asked, uh, Alex Martelli, the editor of the Python Cookbook, which had published my original recipe, a, how do I get this into python? I think it might help

How did we do on your transcript?



Speaker notes

Highlights challenging fragments. Can see what users fix inplace to correct. Star rating for feedback.

SYLLABUS AND CLASS STRUCTURE

17-445/17-645, Summer 2020, 12 units

Tuesday/Wednesday 3-4:20, here on zoom

INSTRUCTORS

Christian Kaestner, Shreyans Sheth

< brief introductions >

COMMUNICATION

Email to us

Announcements through canvas

No fixed office hours, but will stick around after lecture and recitation. Email us for extra meetings.

Welcome to ask questions publicly on Canvas.

Materials on GitHub. Pull requests encouraged!

SOFTWARE ENGINEERING CLASS

- Focused on engineering judgment
- Arguments, tradeoffs, and justification, rather than single correct answer
- "it depends..."
- Practical engagement, building systems, testing, automation
- Strong teamwork component
- Not focused on formal guarantees or machine learning fundamentals (modeling, statistics)

PREREQUISITES

Some software engineering experience required

- version control
- gathering requirements
- software design and modeling
- testing and test automation
- some larger software projects in teams
- see background check quiz on Canvas

No machine-learning knowledge required

- Will cover AI and ML basics this and next week
- If you are familiar with ML/scikit-learn those might be mostly boring... sorry.

In case this is a better fit: We will teach a different version of the class in the Fall that requires some ML experience, but no software engineering experience.

ACTIVE LECTURE

- Case study driven
- Discussion highly encouraged
- Contribute own experience
- Regular active in-class exercises
- In-class presentation
- Discussions over definitions

TEXTBOOK

Building Intelligent Systems: A Guide to Machine Learning Engineering


by Geoff Hulten

<https://www.buildingintelligentsystems.com/>

Most chapters assigned at some point in the
semester

Supplemented with research articles, blog
posts, videos, podcasts, ...

[Electronic version](#) in the library



READINGS AND QUIZZES

- Reading assignments for most lectures
 - Preparing in-class discussions
 - Background material, case descriptions, possibly also podcast, video, wikipedia
 - Complement with own research
- Short and easy online quizzes on readings, due before start of lecture

ASSIGNMENTS

- Series of small to medium-sized individual assignments (mostly in first half)
 - engage with practical challenges
 - design, implement, and automate
 - reason about tradeoffs and justify your decisions
 - written reports and some coding and modeling
- Large team project with multiple milestones (mostly in second half)
 - Build and deploy prediction service
 - Testing in production
 - Monitoring

RECITATIONS

Typically hands on exercises, use tools, analyze cases

Often designed to prepare for assignments

First recitation tomorrow: Apache Kafka

GRADING

- 35% individual assignment
- 30% group project with final presentation
- 20% midterm
- 10% participation
- 5% reading quizzes
- no final exam

PARTICIPATION

- Participation is important
 - Participation in in-class discussions
 - Active participation in recitations
 - Both quality and quantity are important, quality more than quantity
- Participation != Attendance

LATE DAY POLICY

Late work in **individual assignments** will be accepted with a 10% penalty per day, for up to 3 days.

Late work in **group projects** will receive feedback but no credit.

Talk to us (early) for concerns and accommodations.


ACADEMIC HONESTY

See web page

In a nutshell: do not copy, do not lie, do not share or publicly release your solutions

In group work, be honest about contributions of team members, do not cover for others

If you feel overwhelmed or stressed, please come and talk to us (see syllabus for other support opportunities)




ASIDE: AI VS ML

- Artificial intelligence is an umbrella term covering symbolic AI (problem solving, reasoning) as well as machine learning (statistical learning from data)
- This course focuses mostly on *statistical machine learning* and supervised learning (extrapolating from data, inductive reasoning)
- We will cover *symbolic AI* (expert systems, probabilistic reasoning, ...) selectively, often for contrast

PUBLISHING VIDEO RECORDINGS?

With your consent, I'd like to publish recordings publicly

No webcam or chat footage, lightly edited for privacy



Thoughts?

INTRODUCTIONS

Let's go around the "room" for introductions:

- Your (preferred name)
- In two sentences your software engineering background and goals
- In two sentences your data science background, if any, and goals
- One topic you are particularly interested in, if any?



CORRECTNESS AND SPECIFICATIONS

DEDUCTIVE VS. INDUCTIVE REASONING

WHO IS TO BLAME?

```
Algorithms.shortestDistance(g, "Tom", "Anne");  
> ArrayOutOfBoundsException
```

```
Algorithms.shortestDistance(g, "Tom", "Anne");  
> -1
```

WHO IS TO BLAME?

```
class Algorithms {  
    /**  
     * This method finds the shortest distance between to  
     * vertices. It returns -1 if the two nodes are not  
     * connected.  
     */  
    int shortestDistance(...) {...}  
}
```

```
class Algorithms {  
    /**  
     * This method finds the shortest distance between to  
     * vertices. Method is only supported  
     * for connected vertices.  
     */  
    int shortestDistance(...) {...}  
}
```

SYSTEM DECOMPOSITION WITH INTERFACES

```
/*@ requires amount >= 0;
    ensures balance == \old(balance)-amount &&
        \result == balance;
@*/
public int debit(int amount) {
    ...
}
```

(JML specification in Java, pre- and postconditions)

```
/**
 * Calls the <code>read(byte[], int, int)</code> overloaded [..
 * @param buf The buffer to read bytes into
 * @return The value returned from <code>in.read(byte[], int, in
 * @exception IOException If an error occurs
 */
public int read(byte[] buf) throws IOException
{
    return read(buf, 0, buf.length);
}
```

(textual specification with JavaDoc)

CONTRACTS/SPECIFICATIONS

- Contracts describe expected behavior for methods, while hiding the implementation behind
- States method's and caller's responsibilities
- Analogy: legal contract
 - If you pay me this amount on this schedule...
 - I will build the following...
 - Some contracts have remedies for nonperformance
- Invariants must hold before and after loop/method execution
- Defines what it means for implementation to be correct, including exceptional behavior

WHO IS TO BLAME?

```
Math.sqrt(-5);
```

```
> 0
```

BENEFITS OF SPECIFICATIONS

- Exact specification of what should be implemented
- Decompose a system into its parts, develop and test parts independently
- Accurate blame assignments and identification of buggy behavior
- Useful for test generation and as test oracle

SPECIFICATIONS IN MACHINE LEARNING?

```
/**  
    ????  
*/  
String transcribe(File audioFile);
```

SPECIFICATIONS IN MACHINE LEARNING?

```
/**  
    ????  
*/  
List<Product> suggestedPurchases(List<Product> pastPurchases);
```

SPECIFICATIONS IN MACHINE LEARNING?

```
/**  
    ????  
*/  
Boolean predictRecidivism(int age,  
                           List<Crime> priors,  
                           Gender gender,  
                           int timeServed,  
                           ...);
```

SPECIFICATIONS IN MACHINE LEARNING?

- Usually clear specifications do not exist -- we use machine learning exactly because we do not know the specifications
- Can define correctness for some data, but not general rules; sometimes can only determine correctness after the fact
- Learning for tasks for which we cannot write specifications
 - Too complex
 - Rules unknown
- AI will learn rules/specifications, often not in a human-readable form, but are those the right ones?
- Often *goals* used instead --> maximize a specific objective



(Daniel Miessler, CC SA 2.0)

DEDUCTIVE REASONING

- Combining logical statements following agreed upon rules to form new statements
- Proving theorems from axioms
- From general to the particular
- *mathy reasoning, eg. proof that π is irrational*
- Formal methods, classic rule-based AI systems, expert systems

INDUCTIVE REASONING

- Constructing axioms from observations
- Strong evidence suggests a rule
- From particular to the general
- *sciency reasoning, eg. finding laws of nature*
- Most modern machine learning systems, statistical learning

RESULTING SHIFT IN DESIGN THINKING?

From deductive reasoning to inductive reasoning...

From clear specifications to goals...

From guarantees to best effort...

What does this mean for software engineering?

For decomposing software systems?

For correctness of AI-enabled systems?

For safety?

For design, implementation, testing, deployment, operations?

A TOUCH OF REALISM

While it is possible to formally specify programs and prove them correct, this is rarely ever done.

In practice, specifications are often textual, local, weak, vague, or ambiguous, if they exist at all. Some informal requirements and some tests might be the only specifications available.

Software engineers have long development methods to deal with uncertainty, missing specifications, and unreliable components.

AI may raise the stakes, but the problem and solutions are not entirely new.

SURVEY TIME: BACKGROUND SURVEY

Survey helps us to tailor class and form teams (ungraded quiz on Canvas)



SUMMARY

- *Data scientists and software engineers* have different goals and focuses
 - Building systems requires both
 - Various qualities are relevant, beyond just accuracy
- Inductive reasoning and lack of specifications

